

CS-AML

Product Requirements Document

PRD v0.1

Civil Society Financial Intelligence & AML Investigation Platform

Status: Product baseline for MVP planning and engineering handoff

Version: 0.1 | Date: 7 October 2026

Product axiom

The platform SHALL help civil-society investigators produce lawful, evidence-based, reproducible, and actionable financial intelligence while preserving analytical uncertainty and preventing inference from being presented as fact.

Document Control

Field	Value
Document	CS-AML Product Requirements Document (PRD)
Version	0.1
Status	Baseline / implementation planning
Audience	Product, engineering, investigation leads, security/privacy reviewers, QA, governance
Upstream specifications	CS-AML Framework; Goals & Non-Goals; Typology Catalogue; Investigation Methodology; Data Model Specification; Control Implementation Guide; Technology Architecture; Product & Feature Specification
Normative intent	Defines product outcomes and MVP requirements. Detailed data semantics, controls, and architecture remain governed by their respective specifications.

1. Executive Summary

CS-AML is a civil-society financial intelligence and AML investigation platform. It is designed for organisations that investigate corruption, fraud, environmental crime, illicit financial flows, procurement abuse, beneficial-ownership concealment, sanctions/PEP exposure, and related public-interest issues without possessing privileged core-banking access.

The product does not attempt to replicate a bank transaction-monitoring system. Its primary job is to turn lawful public-source and legitimately obtained information into a structured, reviewable intelligence process: Source → Evidence → Fact → Indicator → Hypothesis → Assessment → Intelligence Product.

MVP decision

MVP 0.1 focuses on the complete investigation loop for one organisation: open a case, preserve sources/evidence, resolve entities, map relationships/assets/events/value flows, test typologies and hypotheses, draft an assessment, peer-review it, and generate a controlled intelligence product. Cross-organisation federation, advanced network science, crypto analytics, and autonomous AI are explicitly deferred.

2. Product Problem

2.1 Problem statement

Civil-society investigations frequently begin with fragmented documents, public registries, spreadsheets, media reporting, procurement records, court decisions, tips, and analyst notes. Existing tools may support document search, graph analysis, or case management, but they rarely enforce an end-to-end discipline that preserves provenance, distinguishes fact from inference, supports competing hypotheses, and produces auditable financial-intelligence products.

2.2 User pain points

- Investigators lose time reconstructing where a claim came from and whether a source was preserved.
- Entity duplication creates false relationships and inconsistent case histories.
- Graphs can visually overstate certainty because an edge may look equally “real” whether observed, inferred, or disputed.
- Value flows are often reconstructed from contracts, assets, and ownership events rather than direct bank transactions, yet existing tools do not always preserve that distinction.
- Hypotheses, counter-evidence, intelligence gaps, and confidence judgements are commonly kept in prose rather than structured form.
- Peer review and dissemination approval are often external to the investigative record.
- Cross-case reuse of entities and evidence is difficult without sacrificing compartmentalisation.
- High-risk personal data and protected-source information require stronger controls than generic research tools provide.

3. Product Vision and Positioning

3.1 Vision

Enable civil-society investigators to build defensible financial-intelligence assessments from lawful information with the same discipline expected of professional intelligence work, without pretending to possess investigative powers or financial-system access that they do not have.

3.2 Product category

CS-AML is best positioned as a Civil Society Financial Intelligence & Investigation Platform. It combines case management, evidence/provenance, entity and asset analysis, graph investigation, follow-the-value reconstruction, typology analysis, hypothesis testing, and controlled intelligence-product generation.

3.3 Differentiation

Differentiator	Required product behaviour
Evidence-first	Every material finding can be traced to

Differentiator	Required product behaviour
	source/evidence.
Analytical separation	Fact, indicator, hypothesis, assessment, and allegation remain distinct objects/states.
Follow-the-value	Supports direct, documented, reconstructed, and hypothetical flows without conflating them.
Reversible identity decisions	Entity resolution and merges are evidence-based, reviewable, and reversible.
Civil-society safety	Data minimisation, source protection, compartmentalisation, and controlled dissemination are native.
Intelligence discipline	Alternative hypotheses, disconfirming evidence, confidence, and gaps are part of the workflow—not optional notes.

4. Product Goals and Non-Goals

4.1 Goals

1. Reduce the time required to turn fragmented source material into a structured investigation.
2. Increase the proportion of material findings with complete provenance and evidence lineage.
3. Enable investigators to reconstruct ownership, control, assets, events, and value flows without privileged bank access.
4. Reduce analytical errors caused by entity confusion, unsupported graph edges, confirmation bias, or lost context.
5. Make high-impact intelligence products independently reviewable and reproducible.
6. Provide a secure, role-aware workspace suitable for sensitive civil-society investigations.
7. Provide an architecture that can integrate mature external tools rather than rebuilding generic capabilities unnecessarily.

4.2 Non-goals for MVP 0.1

- Bank-grade real-time transaction monitoring or regulatory STR/SAR submission automation.
- Automatic determination that a person committed money laundering or another offence.
- Autonomous publication, referral, entity merge, or adverse assessment by AI.
- Covert collection, hacking, credential theft, interception, or access to restricted financial systems.
- Full blockchain analytics engine.
- Cross-organisation federation or shared multi-tenant intelligence network.

- Advanced predictive risk scoring presented as an authoritative decision.

5. Target Users and Jobs-to-be-Done

Role	Primary job	Core need
Investigator / Analyst	Build and test a financial-intelligence case	I need to collect, structure, connect, test, and explain information without losing provenance or uncertainty.
Investigation Lead / Case Owner	Control scope, risk, priorities, and decisions	I need to know what is in scope, what remains unknown, and whether the case is ready to advance.
Peer Reviewer / Red Team	Challenge findings independently	I need to trace claims to evidence, inspect alternatives, and identify unsupported leaps.
Evidence Custodian / Data Steward	Protect evidence and data quality	I need to preserve originals, lineage, identifiers, merge history, retention, and integrity.
Security / Privacy Reviewer	Control sensitive data and exposure	I need visibility into classification, access, source protection, export, and retention decisions.
Product Approver / Editor	Approve an intelligence product for external use	I need evidence that identity, uncertainty, legal/privacy risk, and review requirements have been satisfied.
System Administrator	Operate the platform safely	I need secure identity, permissions, backup, monitoring, auditability, and controlled configuration.

6. Core Product Principles

ID	Principle
P-01	Case is context; Entity and Evidence are reusable truth-bearing objects.
P-02	Technology SHALL preserve analytical uncertainty rather than erase it.
P-03	A graph edge is an analytical object with provenance, not merely a visual connection.

ID	Principle
P-04	Direct, documented, reconstructed, and hypothetical value flows SHALL remain distinguishable everywhere.
P-05	Automation MAY suggest; a human SHALL own material analytical and dissemination decisions.
P-06	Sensitive data collection SHALL be necessary, proportionate, and purpose-bound.
P-07	High-impact claims SHALL be independently reviewable from the system of record.
P-08	Derived search/graph/AI indexes are rebuildable projections; canonical evidence and records are authoritative.

7. End-to-End User Journey

1. Intake / lead registration
2. Triage and legitimacy review
3. Investigation Charter
4. Collection plan
5. Source and evidence registration
6. Entity resolution
7. Relationship / ownership / asset mapping
8. Timeline and value-flow reconstruction
9. Typology / indicator mapping
10. Hypothesis testing and intelligence gaps
11. Assessment and confidence
12. Peer review / red-team review
13. Intelligence product generation
14. Dissemination approval / referral / publication
15. Closure / monitoring / reopen

Lifecycle requirement

The UI SHALL make the current investigation stage and unmet gate requirements visible. The

product SHALL NOT silently imply that completion of a form equals analytical validation.

8. MVP 0.1 Scope

8.1 MVP definition

MVP 0.1 is successful when a small civil-society investigation team can complete one sensitive investigation from case opening through reviewed intelligence product using CS-AML as the system of record for the analytical chain.

8.2 MVP capability groups

ID	Capability	MVP content
C1	Case & workflow	Case register, Charter, G0–G6 gates, tasks, activity history
C2	Sources & evidence	Source register, upload, originals, hashes, extracts, derivative lineage, source/credibility ratings
C3	Entities & relationships	Entity registry, aliases, candidate matching, reversible merge/unmerge, first-class relationships, ownership/control
C4	Assets, events & timeline	Asset records, events, timeline view
C5	Value-flow	Flow records, multi-leg flows, classification, visualisation, unknown/range handling
C6	Typology & indicators	Catalogue browser, indicator capture, typology worksheet
C7	Hypothesis & assessment	Hypothesis matrix, support/contradiction, gaps, assessment, confidence, disconfirming-search record
C8	Search & graph	Full-text search, object/entity search, graph exploration
C9	Review & product	Peer review, intelligence-product templates, versioning, evidence index

ID	Capability	MVP content
C10	Dissemination	Approval, secure export, referral package, sharing log
C11	Administration & security	Controlled vocabularies, retention, OIDC/MFA, object-level access, protected-source compartment, immutable audit, backup

8.3 Deferred capability groups

Phase	Deferred capabilities
Phase 2	OCR/table extraction; entity candidate extraction; enrichment connectors; beneficial-ownership chain calculation; path finding; saved searches; OpenAleph/OpenSanctions/FollowTheMoney integrations; AI-assisted extraction/summarisation; operational monitoring.
Phase 3	Network metrics/community detection; partner federation; Flowintel/GraphSense integrations; advanced AI analytical assistant; cross-organisation intelligence exchange.

9. Functional Requirements by Epic

EPIC-01 Case & Workflow

Req ID	Requirement	Product behaviour	MVP acceptance
FR-CASE-01	Create and manage case	Analyst can create a case with owner, purpose, investigation question, scope, sensitivity, jurisdiction, period, and status.	Case record created; mandatory Charter fields validated; audit event created.
FR-CASE-02	Investigation Charter	Case Owner can define included/excluded scope, initial subjects, hypotheses, collection constraints, risks, and intended outputs.	Charter is versioned; scope change requires reason and recorded approval when risk increases.

Req ID	Requirement	Product behaviour	MVP acceptance
FR-CASE-03	Lifecycle gates	System enforces G0–G6 gate records for configured high-risk actions.	Controlled action cannot be marked approved without approver, time, decision, and conditions.
FR-CASE-04	Tasks and assignments	Users can assign investigative tasks with due dates and case linkage.	Task appears in case activity and user queue; completion recorded.

EPIC-02 Source & Evidence

Req ID	Requirement	Product behaviour	MVP acceptance
FR-EVD-01	Source registration	Register source origin, type, URL/location, dates, collector, access/legal note, reliability.	Any material evidence can be traced to a Source ID.
FR-EVD-02	Original evidence preservation	Upload/store original files separately from working derivatives.	Original cannot be silently overwritten; file metadata and hash retained.
FR-EVD-03	Evidence extracts	Create page/section/quote/table/image extracts with exact parent link.	Reviewer can navigate from extract to parent evidence and location.
FR-EVD-04	Derivative lineage	OCR, translation, crop, parsed data, or analyst dataset records parent and transformation metadata.	No derivative can be presented as an original.
FR-EVD-05	Source & information evaluation	Record source reliability separately from information credibility.	UI displays both dimensions independently.

EPIC-03 Entity & Relationship

Req ID	Requirement	Product behaviour	MVP acceptance
FR-ENT-01	Entity registry	Create reusable Person, Organization, Address, Account, Wallet, Asset-related and other supported entity objects.	Stable ID; aliases/identifiers; case links without entity duplication.

Req ID	Requirement	Product behaviour	MVP acceptance
FR-ENT-02	Candidate matching	System surfaces possible duplicate entities using deterministic/fuzzy signals.	Candidate suggestion includes reasons; no automatic merge.
FR-ENT-03	Merge/unmerge	Reviewer-authorized entity merges preserve decision rationale, evidence, aliases, identifiers, and history.	Merge is reversible and audit logged.
FR-REL-01	Relationship records	Create relationship with typed endpoints, evidence, valid time, confidence, status.	Graph edge opens underlying relationship record and evidence.
FR-REL-02	Ownership/control	Represent legal ownership, beneficial ownership, control, use, and association separately.	UI cannot collapse these concepts into one generic “owns” state.

EPIC-04 Assets, Events & Timeline

Req ID	Requirement	Product behaviour	MVP acceptance
FR-AST-01	Asset registry	Register property, vehicle, company share, vessel, crypto asset, or other assets with ownership/control/use distinctions.	Asset has stable ID, evidence, dates, and status.
FR-TIM-01	Event record	Create dated events linked to entities, assets, places, sources/evidence.	Date precision is explicit when exact date is unknown.
FR-TIM-02	Timeline view	Render case/entity timeline from canonical events.	Every timeline item can be opened to supporting record/evidence.

EPIC-05 Follow-the-Value

Req ID	Requirement	Product behaviour	MVP acceptance
FR-VAL-01	ValueFlow record	Record origin, destination, mechanism, date, amount/range/unknown,	Flow class is mandatory: DIRECT, DOCUMENTED, RECONSTRUCTED,

Req ID	Requirement	Product behaviour	MVP acceptance
		currency, evidence, confidence, and flow class.	HYPOTHETICAL.
FR-VAL-02	Multi-leg flow	Link multiple legs into a coherent value-flow chain without losing leg-level evidence.	Each leg retains its own class/evidence; chain does not upgrade certainty.
FR-VAL-03	Flow visualisation	Visualise value flow with visually distinct certainty classes.	Reviewer can distinguish direct vs reconstructed/hypothetical at first glance and via accessible labels.

EPIC-06 Typology & Hypothesis

Req ID	Requirement	Product behaviour	MVP acceptance
FR-TYP-01	Typology catalogue	Browse CS-AML typologies and indicators in-version.	Case records typology catalogue version used.
FR-TYP-02	Typology match worksheet	Map observed indicators/counter-indicators and assessment level.	System prevents typology match from being labelled as proof.
FR-HYP-01	Hypothesis workspace	Create competing hypotheses and link supporting, contradicting, neutral evidence.	At least one alternative explanation can be recorded; status changes audited.
FR-HYP-02	Intelligence gaps	Record unknowns that could materially change judgement.	Gaps appear in assessment and review views.
FR-ASM-01	Assessment & confidence	Create judgement with confidence, basis, assumptions, alternatives, and gaps.	Assessment cannot be final without evidence basis and confidence rationale.

EPIC-07 Search & Graph

Req ID	Requirement	Product behaviour	MVP acceptance
FR-SCH-01	Full-text search	Search permitted	Results respect object-level

Req ID	Requirement	Product behaviour	MVP acceptance
		evidence/documents and analytical records.	access and classification.
FR-SCH-02	Structured search	Filter by entity type, identifiers, dates, relationship type, case, source, status.	Results link to canonical records.
FR-GRF-01	Graph exploration	Explore entities and evidence-backed relationships interactively.	Graph never creates canonical relationships by display alone; inferred/suggested items visibly distinguished.

EPIC-08 Review, Product & Dissemination

Req ID	Requirement	Product behaviour	MVP acceptance
FR-REV-01	Peer review	Independent reviewer records findings, required changes, and approval/rejection.	Author cannot satisfy mandatory independent review alone.
FR-PRD-01	Intelligence product	Generate product from approved assessment, key facts, evidence index, alternatives, gaps, confidence, handling classification.	Output preserves analytical qualifiers and stable evidence references.
FR-PRD-02	Versioning/corrections	Products are versioned; corrections/supersession do not erase prior issued version.	Change history and reason are visible.
FR-DIS-01	Dissemination approval	External dissemination requires authorised decision and handling rules.	Export/referral action creates approval and audit record.
FR-DIS-02	Secure export/referral	Generate controlled package for legitimate recipient.	Package includes classification/handling and evidence index appropriate to recipient.

EPIC-09 Security, Privacy & Operations

Req ID	Requirement	Product behaviour	MVP acceptance
FR-SEC-01	Authentication	OIDC with MFA for privileged and sensitive access.	Failed/privileged auth events logged; session security enforced.
FR-SEC-02	Access control	Role + need-to-know/object-level access.	Unauthorised case/entity/evidence is not discoverable through search, graph, API, export, or logs.
FR-SEC-03	Protected-source compartment	Protected source identity stored separately with more restrictive access.	Routine analyst can use sanitised source reference without viewing identity.
FR-AUD-01	Audit trail	Record material create/update/delete/merge/review/export/dissemination actions.	Audit record is append-only to normal users and queryable for review.
FR-OPS-01	Backup/restore	Back up canonical DB and evidence; test restoration.	Documented restoration test successfully recovers selected case/evidence.

10. UX / Information Architecture Requirements

10.1 Primary navigation

- Home / work queue
- Cases
- Entities
- Evidence & Sources
- Graph
- Search
- Typologies
- Reviews
- Products / Dissemination
- Administration

10.2 Case workspace

Surface	Minimum content
Overview	Question, scope, owner, sensitivity, status, gates,

Surface	Minimum content
	critical gaps
Evidence	Sources, evidence, extracts, derivatives, provenance
Entities	Subjects, organisations, accounts/wallets, addresses, identifiers
Relationships / Graph	Evidence-backed network with confidence/status filters
Assets	Ownership/control/use, valuation where present
Timeline	Events with evidence and precision
Value Flow	Flow chain with certainty class
Typologies	Indicators/counter-indicators and catalogue version
Hypotheses	Competing explanations and evidence matrix
Assessment	Judgement, confidence, basis, gaps, alternatives
Review	Peer/red-team comments, disposition, sign-off
Product	Draft/final intelligence products, exports, sharing log

10.3 UX safety requirements

- Fact, inference, hypothesis, and assessment SHALL use distinct labels and visual treatment.
- Reconstructed and hypothetical flows SHALL never use the same visual style as direct flows.
- Graph and search suggestions SHALL be labelled as suggestions until promoted by a human into canonical records.
- Destructive or high-impact actions (merge, disseminate, delete/retention disposal) SHALL require explicit confirmation and role checks.
- Protected-source identities SHALL never appear in global search snippets, generic logs, or exported reports unless specifically authorised.

11. Non-Functional Requirements

ID	Area	Requirement
NFR-01 Security	Encryption in transit; protected storage for restricted evidence;	

ID	Area	Requirement
	strong authentication; least privilege; secure secrets management.	
NFR-02 Privacy	Purpose binding, minimisation, classification, retention/disposition, access logging for high-sensitivity material.	
NFR-03 Auditability	Material analytical and dissemination changes reconstructable from audit/version history.	
NFR-04 Availability	MVP target: reliable single-organisation service with documented maintenance; restore capability prioritised over high availability.	
NFR-05 Backup/DR	Automated backups; encrypted backup copies; tested restore; RPO/RTO documented before production deployment.	
NFR-06 Performance	Common case/entity/search pages responsive for expected NGO-scale workloads; long analytics handled asynchronously.	
NFR-07 Accessibility	Keyboard-accessible core workflow, readable labels, non-colour-only certainty/status encoding.	
NFR-08 Interoperability	Stable IDs; API/export formats; no requirement that graph/search projections become system of record.	
NFR-09 Explainability	Any material automated score/suggestion exposes contributing signals/version/limitations where applicable.	
NFR-10 Maintainability	Schema migrations, versioned vocabularies, configuration history, environment separation, automated tests.	

12. Data Requirements

12.1 Minimum canonical objects

Case, Source, EvidenceItem, EvidenceExtract, Fact, Entity, Relationship, Asset, Event, ValueFlow, Indicator, TypologyMatch, Hypothesis, IntelligenceGap, Assessment, IntelligenceProduct, Review, Dissemination, AuditEvent.

12.2 Mandatory analytical chain

Canonical chain

SOURCE → EVIDENCE → CLAIM/FACT → INDICATOR → HYPOTHESIS → ASSESSMENT → INTELLIGENCE PRODUCT

12.3 Data integrity constraints

- Material Relationship objects require evidence reference or explicit analytical status explaining why evidence is unavailable.
- Merge decisions require match rationale, supporting/conflicting attributes, reviewer where configured, and reversibility.
- ValueFlow requires flow_class and may record amount as exact, range, or unknown.
- Assessment requires confidence and basis; high-impact assessment requires review record.
- Derivatives require parent evidence; original evidence cannot be silently replaced.
- All canonical objects use stable IDs and version/audit semantics defined by the Data Model Specification.

13. Integration Strategy

System	Phase	Purpose	Product rule
OpenAleph	Phase 2	Document ingestion/search/investigative workspace reference or connector	Do not duplicate mature document-processing capabilities if integration satisfies provenance/access requirements.
FollowTheMoney	Phase 2	Ontology/interoperability mapping	Map where compatible; preserve CS-AML-specific Evidence/Hypothesis/Assessment semantics.
OpenSanctions	Phase 2	PEP/sanctions/entity matching	Screening is contextual signal, not proof of wrongdoing.
Neo4j / Memgraph	Phase 2	Derived graph projection /	Canonical Relationship

System	Phase	Purpose	Product rule
		analytics	remains source of truth.
Flowintel	Phase 3	Case/task interoperability	Use only if workflow interoperability provides value without duplicating CS-AML analytical state.
GraphSense	Phase 3	Cryptoasset analytics	Connector rather than building chain analytics.
OIDC / Keycloak-compatible	MVP	Authentication/identity	Use standards-based IAM; MFA and role claims supported.

14. AI and Automation Requirements

14.1 Permitted assistive uses

- OCR/text extraction assistance
- Translation assistance
- Entity candidate extraction
- Summarisation with source links
- Search/query assistance
- Suggested relationship or typology candidates
- Drafting assistance for analyst-controlled products

14.2 Prohibited autonomous decisions

- Determining guilt or criminality
- Publishing/referring a case
- Merging entities without human decision
- Converting a suggestion/inference into a material fact
- Revealing protected-source identity
- Overriding access controls or case compartmentalisation

14.3 Verification

AI output used materially SHALL remain a derived analytical artefact until verified against underlying evidence. The system SHOULD retain model/tool version and prompt/context provenance when needed for reproducibility without exposing unnecessary sensitive content.

15. Success Metrics

Metric	Target / interpretation
Provenance coverage	≥95% of material facts in pilot products trace to registered evidence.
Review completeness	100% of high-impact products have independent review before dissemination.
Entity resolution quality	All merged high-impact entities have recorded rationale; correction/unmerge rate monitored.
Analytical discipline	≥90% of pilot assessments record alternatives and material intelligence gaps.
Workflow adoption	Pilot investigators complete an end-to-end case without maintaining a parallel “shadow” system for core analytical state.
Safety	0 unauthorised disclosures of protected-source identity during pilot.
Recovery	Successful tested restore of a selected case and evidence set before production launch.
Usefulness	Pilot users rate evidence traceability, graph/value-flow understanding, and review workflow as materially better than baseline process.

16. MVP Acceptance / Release Gate

MVP 0.1 release gate

The release is not “done” because all screens exist. It is done when a representative case can be completed end-to-end with provenance, access controls, analytical separation, peer review, export controls, and restore/audit evidence functioning together.

8. Create a case and approved Investigation Charter.
9. Register at least three different source/evidence types and preserve originals.
10. Create evidence extracts with lineage.
11. Create and resolve duplicate entity candidates, including one reversible merge.
12. Create evidence-backed relationships and render them in graph view.
13. Record asset/event/timeline data.
14. Build a multi-leg value-flow chain containing at least two different flow classes.

15. Apply one typology worksheet with counter-indicators.
16. Test at least two competing hypotheses and record an intelligence gap.
17. Create an assessment with confidence and disconfirming-search record.
18. Complete independent peer review.
19. Generate an intelligence product with evidence index and handling classification.
20. Approve and export a controlled referral/share package.
21. Confirm access-control isolation for a restricted case and protected-source identity.
22. Restore the pilot case and evidence from backup and verify audit continuity.

17. Product Risks and Mitigations

ID	Risk	Failure mode	Mitigation
R-01	Scope explosion	Trying to implement all 87 catalogue features at once.	Freeze MVP capability list; phase integrations and advanced analytics.
R-02	Graph overclaim	Users infer wrongdoing from visual proximity.	Evidence-backed edges, certainty labels, legends, relationship status and training.
R-03	Entity false merge	Common names/addresses create false networks.	Candidate state, reviewer decision, merge rationale, reversible merges.
R-04	Sensitive data exposure	Protected source or private data leaks through search/export/logs.	Compartmentalisation, object permissions, sanitised references, export approval, security tests.
R-05	Analytical bureaucracy	Framework creates forms without improving investigation.	Measure outcomes/usefulness; keep gates risk-based; reduce low-value mandatory fields.
R-06	AI authority creep	Users treat AI suggestion as verified fact.	Derived-state labels, verification gate, no autonomous merge/dissemination/adverse judgement.
R-07	Tool duplication	Rebuild OCR, screening, crypto, or document search poorly.	Build-vs-integrate reviews and connectors to mature tools.
R-08	Legal/jurisdiction variance	Lawful collection and publication differ by jurisdiction.	Configurable policy; legal review gate; do not hard-code one jurisdiction as universal law.

18. Delivery Plan

18.1 Suggested releases

Release	Outcome	Scope
0.1 Alpha	Core analytical chain works	Case, evidence, entity/relationship, assets/events, value-flow, typology, hypothesis, assessment, basic graph/search, auth/audit.
0.1 Beta	Safe review and dissemination	Peer review, product generation, classification, secure export/referral, backup/restore, pilot security/privacy tests.
0.1 Production	Operational baseline	Pilot feedback resolved, conformance evidence, operational documentation, monitoring, incident/restore procedures.
0.2	Investigation acceleration	Document processing, enrichment, BO chains, path finding, saved searches, first connectors, AI assist.
0.3	Network intelligence	Cross-case analytics, advanced graph metrics, federation/partner workflows, crypto connector, controlled advanced AI.

18.2 Suggested engineering epics

- E1 Platform foundation & IAM
- E2 Canonical data model & API
- E3 Case/workflow & gates
- E4 Evidence/provenance storage
- E5 Entity resolution & relationships
- E6 Assets/events/timeline
- E7 Value-flow engine
- E8 Typology/hypothesis/assessment
- E9 Search & graph projection
- E10 Review/product/dissemination
- E11 Audit/security/privacy controls
- E12 Backup/restore/ops & pilot readiness

19. Dependencies

Dependency	Why it matters	Owner / decision
Canonical Data Model v0.1	Objects, statuses, IDs, lineage and validation drive API/database.	Architecture/Data
Investigation Methodology v0.1	Defines lifecycle and gates.	Investigation/Product
Control Implementation Guide v0.1	Defines security, privacy, review and audit controls.	Governance/Security
Typology Catalogue v0.1	Provides versioned typology content and indicators.	Methodology/AML
Technology Architecture v0.1	Constrains canonical vs derived stores, security, deployment and AI.	Architecture
Legal/privacy policy	Defines jurisdiction-specific lawful collection, retention, publication and sharing.	Governance/Legal
Pilot cases/users	Needed for usability, conformance and outcome testing.	Product/Investigation

20. Open Product Decisions

ID	Decision	Recommended baseline
D-01	Single-organisation first or multi-tenant from day one?	Recommend single-organisation logical deployment for MVP; design IDs/access boundaries to permit later tenancy.
D-02	Native document processing vs OpenAleph-first?	Recommend minimal native upload/text metadata in MVP; evaluate OpenAleph connector for Phase 2.
D-03	Graph store at MVP?	Recommend relational canonical relationships and lightweight projection first; dedicated graph DB only when path/network workloads justify it.

ID	Decision	Recommended baseline
D-04	How strict are lifecycle gates?	Make gate enforcement configurable by risk profile; never make high-impact dissemination optional.
D-05	Which jurisdictions first?	Keep core implementation-neutral; create local policy pack for Indonesia as an initial deployment option.
D-06	Offline/air-gapped support?	Decide based on threat model of pilot organisations; architecture should not assume permanent cloud connectivity.
D-07	Public/open-source release model?	Decide licensing, governance, security disclosure, and public roadmap before external release.

21. Definition of Done

A feature is done only when its user-visible behaviour, canonical data state, permissions, audit events, tests, error handling, migration implications, and documentation are complete. For security/privacy-sensitive features, threat/abuse cases and negative tests are part of the definition of done.

Dimension	Done when
Product	User story and acceptance criteria satisfied.
Data	Canonical objects validate and preserve provenance/version semantics.
Security	Authorisation and sensitive-data rules pass positive and negative tests.
Audit	Material actions create expected audit events.
UX	Status/uncertainty is legible and accessible.
QA	Automated tests plus representative end-to-end scenario pass.
Operations	Backup/restore/monitoring impact documented where

Dimension	Done when
	relevant.
Documentation	User/admin/API documentation updated.

22. Traceability to CS-AML Specifications

PRD area	Upstream authority
Goals/non-goals	CS-AML Framework Goals & Non-Goals
Lifecycle / gates	CS-AML Investigation Methodology
Objects / states / lineage	CS-AML Data Model Specification
Typologies / indicators	CS-AML Typology Catalogue
Security/privacy/audit/review controls	CS-AML Control Implementation Guide
System architecture / canonical vs derived / AI	CS-AML Technology Architecture
Feature IDs / prioritisation master list	CS-AML Product & Feature Specification

Annex A — MVP P0 Feature IDs

The Product & Feature Specification remains the authoritative detailed feature catalogue. The following list identifies the P0 baseline expected to be represented in MVP planning; final sprint sequencing is an engineering/product decision.

P0 Feature	P0 Feature	P0 Feature
F-CASE-001	F-CASE-002	F-CASE-003
F-CASE-004	F-CASE-005	F-EVD-001
F-EVD-002	F-EVD-003	F-EVD-004
F-EVD-005	F-EVD-006	F-DOC-001
F-ENT-001	F-ENT-002	F-ENT-003
F-ENT-004	F-ENT-005	F-REL-001

P0 Feature	P0 Feature	P0 Feature
F-REL-002	F-REL-003	F-AST-001
F-TIM-001	F-TIM-002	F-VAL-001
F-VAL-002	F-VAL-003	F-VAL-004
F-TYP-001	F-TYP-002	F-TYP-003
F-HYP-001	F-HYP-002	F-HYP-003
F-ASM-001	F-ASM-002	F-ASM-003
F-SCH-001	F-SCH-002	F-GRF-001
F-PRD-001	F-REV-001	F-PRD-002
F-PRD-003	F-DIS-001	F-DIS-002
F-DIS-003	F-DIS-004	F-ADM-001
F-ADM-003	F-SEC-001	F-SEC-002
F-SEC-003	F-AUD-001	F-OPS-001

Annex B — Representative End-to-End Pilot Scenario

A pilot case should use synthetic or legally cleared data and include: one public procurement contract, two organisations with partially overlapping identifiers, one person with an alias, one asset, one time-bounded ownership relationship, one reconstructed value-flow leg, one direct/documented leg, one typology with a counter-indicator, two competing hypotheses, one protected source reference, one peer-review challenge, and one controlled referral/export. The test succeeds only if a reviewer can reproduce the final assessment without relying on undocumented analyst memory.