

CS-AML

Information Architecture Specification

Version 0.1.4

Approved Internal Specification Baseline

2026-10-10 · tag v0.1.4-spec

Civil Society Anti-Money Laundering & Financial Intelligence Framework

Generated from the canonical Markdown source. The Markdown file in the repository remains authoritative; change tags such as `[v0.1.x · ...]` are part of the traceability record.

Contents

1. Purpose and Scope	4
1.1 In scope	5
1.2 Out of scope	5
2. IA Design Principles	5
3. Information Domains	6
4. Canonical Object Hierarchy	6
4.1 Canonical page eligibility	7
5. Global Navigation Model	7
6. Role-Based Entry Points	7
7. Case Workspace Information Architecture	8
7.1 Case context persistence	8
8. Canonical Object Page Pattern	8
9. Entity Information Architecture	9
9.1 Entity taxonomy	9
9.2 Entity page sections	9
9.3 Entity status labels	9
10. Relationship, Ownership and Control Architecture	10
11. Asset and Event Architecture	10
11.1 Asset information	10
11.2 Event taxonomy	10
12. Value-Flow Information Architecture	11
12.1 ValueFlow page structure	11
12.2 Multi-leg chain structure	11
13. Analytical Reasoning Architecture	11
13.1 Typology information	11
13.2 Hypothesis and assessment grouping	12
14. Intelligence Product and Dissemination Architecture	12
14.1 Product taxonomy	12
14.2 Version structure	12
14.3 Dissemination metadata	12
15. Taxonomy and Controlled Vocabulary	12
15.1 Vocabulary rules	13
16. Labeling System	13
16.1 Label priorities	13
16.2 Naming pattern	13
17. Search and Findability Architecture	13
17.1 Search scopes	13
17.2 Search result architecture	14
17.3 Facets	14
18. Cross-Object Discovery and Backlinks	14
19. Deep Linking and Canonical URLs	14

20. Information Scent and Context Cues	15
21. Permission-Aware Information Architecture	15
22. Derived Views: Graph, Timeline and Dashboards	15
23. Content Lifecycle, Versioning and History	16
24. IA Requirements Catalogue	16
25. IA-to-UX Handoff Contract	16
26. IA Validation and Testing	17
26.1 Validation methods	17
26.2 MVP IA acceptance scenarios	17
27. IA Definition of Done	17
28. Traceability to CS-AML Suite	18
Annex A — Reference Navigation Tree	18
Annex B — Core Label Dictionary	18
Annex C — Information Architecture Change Record	19

CS-AML

Information Architecture Specification

Version 0.1.4

Information organisation, taxonomy, navigation, labeling and findability for the CS-AML Platform

Document status — v0.1.4

Version: 0.1.4 — Approved Internal Specification Baseline (2026-10-10, tag v0.1.4-spec). Supersedes v0.1.3 (2026-10-09, tag v0.1.3-spec). [v0.1.4]

Supersedes: CS-AML Information Architecture Specification v0.1. The DOCX/PDF files in this repository are the unchanged v0.1 baseline (legacy); this Markdown file is the canonical source.

Validation: approved by the product owner as the internal specification baseline on 2026-10-08 (v0.1.1) and 2026-10-09 (v0.1.3) and 2026-10-10 (v0.1.4; decision register and release gates in CHANGELOG.md). The v0.1.3 and v0.1.4 changes come from change requests raised while implementing increments I5-I7 and the post-MVP follow-ups; this is not an independent audit. Acceptance criteria in this document are targets, not evidence that tests have passed.

CS-AML is not an external standard or certification. References to FATF, Wolfsberg, PPATK, UNODC or other bodies do not imply their endorsement.

Changes in 0.1.1: see CHANGELOG.md at the repository root (audit findings A01-A16).

Changes in 0.1.3: change requests CR-I5-01...CR-I7-08 approved by the product owner on 2026-10-09 (CHANGELOG.md, section v0.1.3). Each change is tagged *[v0.1.3 · CR-xx-yy]*.

Changes in 0.1.4: change requests CR-N-01...CR-N-14 approved by the product owner on 2026-10-10 (CHANGELOG.md, section v0.1.4). Each change is tagged *[v0.1.4 · CR-N-xx]*.

Status

Approved Internal Specification Baseline (2026-10-10, tag v0.1.4-spec) — normative information-architecture baseline for MVP 0.1. [v0.1.1 · A01] This document defines how information is organised, named, grouped, navigated, linked, searched and found. Interaction behaviour remains governed by the separate CS-AML UX Specification.

Document	Primary ownership
UX Specification v0.1.1 [v0.1.1 · A01]	Interaction behaviour, task flow, feedback, user-facing safety, accessibility.
This IA Specification	Object hierarchy, taxonomy, labels, navigation model, information scent, findability, cross-object discovery.
Data Model Specification	Canonical object semantics, fields and integrity rules.
SRS / PRD	Product and software requirements that IA must satisfy.
Technical Stack	Implementation constraints; IA remains technology-neutral where possible.

CS-AML Framework Suite • 2026

1. Purpose and Scope

This specification defines the information architecture of the CS-AML platform. Its purpose is to make a complex investigation domain understandable and findable without collapsing analytical distinctions. It converts the canonical data model and product capabilities into a coherent user-facing information system.

IA axiom

Case is context; canonical objects are reusable. The information architecture SHALL organise investigation work around context while avoiding unnecessary duplication or imprisonment of entities, evidence and analytical objects inside a single case.

1.1 In scope

- Information domains and object hierarchy.
- Primary and secondary navigation model.
- Case workspace information structure.
- Canonical object pages and contextual object views.
- Taxonomy and controlled vocabulary presentation.
- Labeling conventions and naming rules.
- Global and case-scoped search architecture.
- Facets, filters, browse paths and information scent.
- Cross-object links, backlinks and related-object discovery.
- Deep links, canonical URLs and context preservation.
- Role-based entry points and visibility rules.
- Content grouping, metadata hierarchy and progressive disclosure from an IA perspective.

1.2 Out of scope

- Micro-interaction behaviour, confirmation dialogs and undo mechanics.
- Visual design system details such as component styling or color tokens.
- Keyboard behaviour and accessibility interaction patterns.
- Backend schema implementation beyond IA-to-data mapping.
- Final analytics algorithms or scoring models.

These are governed by UX, SRS, Data Model, Technology Architecture and implementation specifications.

2. IA Design Principles

IA-P01 Canonical object first

Entity, Evidence, Relationship, Asset, Event, ValueFlow and other reusable objects SHALL have canonical identity independent of any single case.

IA-P02 Case as context

Case pages organise investigation work, permissions, tasks, questions and analytical outputs; they do not redefine canonical truth objects.

IA-P03 Provenance is navigable

Users SHALL be able to move from analytical conclusion to supporting evidence and source through explicit information links.

IA-P04 Uncertainty is structured

Status, confidence, dispute, unknowns and analytical class SHALL be represented as information attributes, not hidden annotations.

IA-P05 Findability without leakage

Search and browse SHALL help authorised users discover relevant objects while not revealing restricted object existence where policy forbids it.

IA-P06 Stable terminology

Core labels SHALL map consistently to canonical objects and controlled vocabularies.

IA-P07 Multiple paths, one object

Users MAY arrive at an object from case, search, graph, timeline or report, but SHALL reach the same canonical record.

IA-P08 Relationships are first-class

Relationships SHALL be independently addressable information objects with evidence and temporal context, not merely visual graph edges.

IA-P09 Derived views are not truth stores

Graph, timeline, search result and dashboard views are projections over canonical data.

IA-P10 Progressive depth

Primary surfaces expose the information needed for decisions; detailed provenance, audit and technical metadata remain available at deeper levels.

3. Information Domains

Domain	Primary objects	Purpose
Governance & Case	Case, InvestigationQuestion, Task, ReviewGate	Purpose, authority, scope and work coordination.
Sources & Evidence	Source, EvidenceItem, EvidenceExtract, DerivedArtifact	Provenance, original material and precise citations.
Identity & Network	Entity, Identifier, Relationship, OwnershipInterest, ControlAssertion	Who/what exists and how objects are connected.
Assets & Events	Asset, Event	What value-bearing objects exist and what happened when.
Value Movement	ValueFlow, ValueFlowLeg	Movement/transformation of economic value.
Analytical Reasoning	Indicator, TypologyMatch, Hypothesis, IntelligenceGap, Assessment	How observations become reasoned judgements.
Products & Dissemination	IntelligenceProduct, Review, Dissemination	How intelligence is reviewed, versioned and shared.
Administration & Assurance	Vocabulary, RetentionPolicy, AuditEvent	How semantics, lifecycle and accountability are governed.


4. Canonical Object Hierarchy

The IA SHALL distinguish canonical objects from contextual associations and derived views. Canonical objects receive stable identifiers and canonical object pages when relevant to users.

```

CASE (context)
├─ Investigation Question / Scope / Tasks / Gates
├─ links to reusable canonical objects
│   └─ SOURCE → EVIDENCE → EXTRACT
│   └─ ENTITY ↔ RELATIONSHIP ↔ ENTITY
│   └─ ENTITY ↔ ASSET

```



4.1 Canonical page eligibility

Object	Canonical page	Case-context view	Rationale
Case	Yes	N/A	Primary investigation context.
Source	Yes	Yes	Reusable provenance source.
EvidenceItem	Yes	Yes	Evidence may support multiple claims/cases.
EvidenceExtract	Yes/anchored	Yes	Precise citation to parent evidence.
Entity	Yes	Yes	Reusable across cases.
Relationship	Yes	Yes	First-class analytical proposition.
Asset	Yes	Yes	Reusable object with attribution.
Event	Yes	Yes	Temporal fact/context.
ValueFlow	Yes	Yes	Reusable value-movement record.
Hypothesis	Case-bound	Yes	Analytical reasoning belongs to investigation context.
Assessment	Case-bound	Yes	Judgement depends on scope/time/evidence snapshot.
IntelligenceProduct	Yes/versioned	Yes	Stable product version and dissemination history.

5. Global Navigation Model

The MVP global navigation SHALL remain small and task-oriented. It should expose major information domains without mirroring backend modules.

Primary destination	Purpose	Typical objects
Home	Role-aware entry point, assigned work, recent authorised activity.	Cases, Tasks, Reviews
Cases	Browse and enter investigation contexts.	Case
Entities	Browse/search canonical people, organisations, accounts, addresses and other entities.	Entity, Identifier
Evidence	Browse authorised sources and evidence outside a single case when policy allows.	Source, EvidenceItem
Analysis	Entry to typologies, hypotheses/assessments and analytical work queues.	Typology, Hypothesis, Assessment
Products	Draft, review and approved intelligence products.	IntelligenceProduct, Review, Dissemination
Search	Global permission-aware discovery across canonical objects.	All indexed objects
Administration	Policy-controlled vocabularies, retention, user/system governance.	Vocabulary, RetentionPolicy

Navigation restraint

Graph, Timeline and Value Flow SHOULD NOT be primary global destinations in MVP. They are analytical views anchored in a case, entity or selected object set. This prevents derived projections from competing with canonical information domains.

6. Role-Based Entry Points

Role	Default entry emphasis	Do not assume
Investigator	Assigned cases, tasks, recently viewed evidence/entities.	Access to all organisational investigations.
Case Owner	Case health, gate status, pending approvals, unresolved risks.	Administrative system privileges.
Reviewer	Review queue, submitted products/assessments, evidence traceability.	Need to browse unrelated cases.
Evidence Custodian	Ingestion queue, integrity status, derivatives, retention/hold.	Authority to alter analytical judgement.
Data Steward	Entity match/merge queue, identity-quality issues.	Automatic visibility into protected sources.
Source Handler	Protected-source workspace and controlled handoff.	Protected identity in global search.
Admin/Auditor	Policy/configuration/audit entry points.	Analytical edit rights.

Roles and classification. Role entry points shape default emphasis only; they do not grant access to any classification level. Visibility is decided by server-side policy against the object's classification level (PUBLIC, INTERNAL, SENSITIVE, RESTRICTED, SOURCE_PROTECTED) plus its access labels. No role by itself grants RESTRICTED access (named-role or case-specific authorization is required), and SOURCE_PROTECTED material stays in the protected-source workspace (Source Handler) and out of general navigation and search. Unknown or missing classification fails closed. [v0.1.1 · A08] Exception decided in v0.1.3: a principal holding the per-case protected-source grant MAY include SOURCE_PROTECTED material in search by an explicit opt-in (include_source_protected); by default, and always without the grant, it produces no result, count, facet or snippet. Protected-source identity never appears in general search (Section 21). [v0.1.3 · CR-15-04]

7. Case Workspace Information Architecture

A case workspace SHALL organise investigation context around a stable case header plus domain views. Exact visual interaction is governed by UX; this section defines information grouping and hierarchy.

Case section	Information responsibility	Primary objects
Overview	Question, purpose, status, owner, classification, scope summary, gate state.	Case, InvestigationQuestion
Work	Tasks, assignments, milestones and case activity.	Task, AuditEvent
Sources & Evidence	Case-linked sources, evidence, extracts and derivatives.	Source, EvidenceItem, EvidenceExtract
Entities & Assets	Case-linked canonical entities, identifiers, relationships and assets.	Entity, Relationship, Asset
Timeline	Chronological projection of relevant events.	Event
Value Flows	Case-linked value movement and multi-leg chains.	ValueFlow, ValueFlowLeg
Analysis	Indicators, typology matches, hypotheses, gaps and assessments.	Indicator, TypologyMatch, Hypothesis, IntelligenceGap, Assessment
Products	Draft/reviewed products, reviews and dissemination records.	IntelligenceProduct, Review, Dissemination

7.1 Case context persistence

When a user opens a canonical entity, evidence item or relationship from within a case, the system SHOULD preserve the originating case context so case-specific associations and analytical relevance can be shown without creating a duplicate object.

8. Canonical Object Page Pattern

Canonical object pages SHOULD use a consistent information hierarchy regardless of object type.

Layer	Contents	Rule
Identity	Canonical name/title, stable ID, object type, status/classification.	Always visible near page header.
Core properties	Object-specific primary attributes.	Prioritise authoritative/current values while retaining conflicts.
Provenance	Source/evidence basis, creation and verification context.	Directly reachable from material assertions.
Relationships	Inbound/outbound links to other canonical objects.	Permission-aware and typed.
Case contexts	Authorised cases in which the object is relevant.	Case association must not imply guilt.
History	Temporal changes, supersession, merge/version history.	Do not overwrite history.
Audit/technical	Audit trail and system metadata where role permits.	Progressively disclosed.

9. Entity Information Architecture

9.1 Entity taxonomy

Top-level entity family	Examples	Label rule
Person	Individual, public official, professional intermediary.	Use "Person", not suspect/target as type.
Organisation	Company, nonprofit, agency, partnership, informal organisation.	Specific subtype may be shown as metadata.
Account/ Instrument	Bank account, payment account, virtual asset wallet.	Avoid implying ownership unless relationship supports it.
Place/Address	Registered address, property address, operational location.	Separate location object from ownership assertion.
Digital Identifier	Email, phone, domain, username/device identifier where lawful.	Show source and verification status.
Other Legal/ Economic Object	Contract, project or instrument when modelled as entity-like object.	Prefer dedicated domain objects where defined.

9.2 Entity page sections

- Identity summary and status.
- Aliases and identifiers with provenance.
- Relationships grouped by relationship family and time.
- Assets and value-flow participation.
- Events/timeline references.
- Evidence and source support.
- Case contexts.
- Merge/unmerge and identity-resolution history for authorised roles.

9.3 Entity status labels

Status	Meaning
Candidate	Possible entity record or unresolved identity.
Probable	Evidence suggests identity but material uncertainty remains.
Confirmed	Identity resolution is sufficiently supported for current analytical purpose.
Disputed	Material conflicting evidence exists.
Unresolved	Available information is insufficient to decide.

Display labels map to Data Model wire values: Confirmed → RESOLVED; Disputed → CONFLICTED; Unresolved → UNRESOLVED; plus MERGED and SPLIT. Candidate and Probable describe a pending or POSSIBLE_MATCH ResolutionDecision between records, not a stored entity state. Entity state changes only through ResolutionDecisions (Data Model v0.1.4 §8.4). [v0.1.1 · ER]

10. Relationship, Ownership and Control Architecture

Relationships are first-class records and SHALL be discoverable from both endpoints as well as case context.

Relationship family	Examples	Required information scent
Governance/role	DIRECTOR_OF, EMPLOYED_BY, REPRESENTED_BY	Role type, organisation, valid time, evidence/status.
Ownership	OWNS, BENEFICIAL_OWNER_OF	Owner/asset, percentage if known, direct/indirect, evidence.
Control	CONTROLS, MANAGES, USES [v0.1.1 · C09]	Control basis distinct from legal ownership.
Commercial	SUPPLIER_TO, CONTRACTED_BY	Counterparty, contract/project context, dates.
Personal/association	RELATED_TO (familial subtype; displayed as "Relative of"), ASSOCIATE_OF [v0.1.1 · C09]	Use carefully; provenance and relevance required.
Transfer/value	PAID_BY, TRANSFERRED_VALUE_TO, ACQUIRED, SOLD_TO [v0.1.1 · C09]	Prefer ValueFlow/Event where movement semantics are material.

Graph rule

A graph edge SHALL resolve to a canonical Relationship, OwnershipInterest, ControlAssertion or ValueFlow object. The graph itself is never the sole record of the connection.

11. Asset and Event Architecture

11.1 Asset information

- Asset identity/type and identifiers.
- Location/jurisdiction where relevant.
- Legal ownership, beneficial ownership, control, use and association as distinct relationships.
- Valuation records with source/date/method where present.
- Events affecting acquisition, transfer, sale, seizure or status.
- Supporting evidence and case contexts.

11.2 Event taxonomy

Event family	Examples	Date semantics
Organisation	incorporation, director change, dissolution	Exact/range/precision must be explicit.
Procurement/contract	award, subcontract, payment milestone	Link project/contract/entity/evidence.
Asset	acquisition, transfer, sale, valuation event	Do not infer causation from temporal proximity.
Legal	court filing, order, sanction, administrative action	Jurisdiction and source are important.
Analytical	case-opened/reviewed may appear in case activity, not factual event timeline by default	Separate investigation activity from subject-world events.

12. Value-Flow Information Architecture

ValueFlow is a canonical analytical object representing movement or transformation of economic value. Its information architecture SHALL preserve evidentiary class.

Class	IA meaning	Required label
DIRECT	Direct transaction evidence supports the movement.	DIRECT
DOCUMENTED	Documented economic relationship/obligation supports the movement but not necessarily transaction-level proof.	DOCUMENTED
RECONSTRUCTED	Analyst reconstructs probable movement from linked facts and events.	RECONSTRUCTED
HYPOTHETICAL	Proposed flow used for hypothesis testing.	HYPOTHETICAL

The required labels are the `flow_class` wire values (UPPER_SNAKE_CASE, derived from the Data Model Annex A registry); translated display labels MAY be shown alongside them but SHALL NOT replace or merge classes. [v0.1.1 · A09]

12.1 ValueFlow page structure

- Origin and destination.
- Mechanism/type.
- Value, currency, range/unknown state and date/precision.
- Flow class and confidence/status.
- Evidence and source basis for each leg.
- Related contract, asset, event and entity objects.
- Case contexts and analytical usage.

12.2 Multi-leg chain structure

A chain is a view over ordered ValueFlowLeg objects. Each leg retains independent evidence, class and confidence. The chain label SHALL not override per-leg semantics.

13. Analytical Reasoning Architecture

```

OBSERVATION / FACT
  ↓
INDICATOR ↔ COUNTER-INDICATOR
  ↓
TYPOLOGY MATCH (consistency, not proof)
  ↓
HYPOTHESES + CONTRADICTIONS + GAPS
  ↓
ASSESSMENT + CONFIDENCE
  ↓
INTELLIGENCE PRODUCT

```

13.1 Typology information

- Catalogue entry ID/version/name.
- Mechanism and objective.
- Observables and indicators.
- Counter-indicators / common false positives.
- Relevant evidence classes and analytical cautions.
- Case-specific TypologyMatch records SHALL remain separate from catalogue definition.

13.2 Hypothesis and assessment grouping

Object	Grouping rule	Cross-links
Hypothesis	Case-bound; grouped by status and analytical question.	Supporting/contradicting evidence, indicators, gaps.
IntelligenceGap	Grouped by open/closed and materiality.	Hypothesis/assessment/tasks.
Assessment	Versioned by case/question; not global truth about an entity.	Hypotheses, evidence basis, confidence, alternatives, product versions.

14. Intelligence Product and Dissemination Architecture

14.1 Product taxonomy

Product type	Primary use
Financial Intelligence Note	Concise assessment for internal or controlled external use.
Entity Profile	Structured identity, relationships, assets and provenance summary.
Asset Profile	Asset ownership/control/use and relevant events/value context.
Network Analysis	Evidence-backed network findings and analytical interpretation.
Referral Package	Structured package for competent authority or trusted recipient.
Case Report	Comprehensive investigation record and conclusions.

14.2 Version structure

Each product SHALL have stable Product ID plus immutable version identity. Current, superseded and retracted versions SHALL be distinguishable. Dissemination records link to the exact version shared, not merely the product family.

14.3 Dissemination metadata

- Recipient / recipient class.
- Purpose.
- Date/time.
- Product version.
- Classification/handling restrictions.
- Included/excluded evidence scope.
- Approver and approval reference.
- Correction/update status where relevant.

15. Taxonomy and Controlled Vocabulary

Taxonomy SHALL distinguish system-defined canonical categories from organisation-defined controlled vocabularies and analyst-entered free text.

Layer	Examples	Governance
Canonical system enum	ValueFlow class, analytical status, product lifecycle state, confidence level (HIGH, MODERATE, LOW, INSUFFICIENT_BASIS), classification level (PUBLIC, INTERNAL, SENSITIVE, RESTRICTED, SOURCE_PROTECTED).	Versioned through software/data model (Data Model Annex A registry); wire values are UPPER_SNAKE_CASE and display labels are separate and translatable; changes require migration review. [v0.1.1 · A08, A09]

Layer	Examples	Governance
Framework catalogue	CS-AML typology IDs and classes.	Versioned framework release.
Controlled vocabulary	Relationship types, asset types, event types, access labels (e.g. purpose, jurisdiction, embargo, legal-review, compartment).	Admin-managed with stable IDs and effective dates. Classification levels are a canonical system enum, not an admin-managed vocabulary. [v0.1.1 · A08]
Reference taxonomy	Jurisdiction lists, currencies, organisation types.	May map to external standards where appropriate.
Free text	Notes, rationale, narrative assessment.	Not used where controlled semantics are required.

15.1 Vocabulary rules

- Labels MAY change, stable IDs SHOULD not.
- Deprecated terms SHALL remain interpretable on historical records.
- Synonyms MAY support search but SHALL not silently merge distinct concepts.
- Relationship direction SHALL be explicit and user-facing labels understandable in both directions where needed.
- Taxonomies SHALL avoid loaded terms such as “criminal”, “dirty”, or “suspicious entity” as structural types unless legally established and contextually required.

16. Labeling System

16.1 Label priorities

Priority	Rule	Example
1	Use domain language that matches framework/data objects.	“Evidence”, “Hypothesis”, “Assessment”.
2	Prefer concrete nouns for destinations and objects.	“Entities” rather than “Network Intelligence Centre”.
3	Use verbs for actions, not destinations.	“Add evidence”, “Create relationship”.
4	Avoid accusatory labels for neutral states.	“Case subject” only as contextual role, not entity type.
5	Show uncertainty in labels when material.	“Candidate match”, “Reconstructed flow”.

16.2 Naming pattern

Object type + canonical label + status/context

Example:

Company · PT Example Nusantara · Confirmed

Value Flow · Company A → Company B · RECONSTRUCTED

Assessment · Procurement ownership question · Moderate confidence

17. Search and Findability Architecture

17.1 Search scopes

Scope	Purpose	Default boundary
Global	Find authorised canonical objects across platform.	User permissions and object classification.
Case	Find objects associated with current case.	Case membership plus object permission.
Evidence/ document	Find text/citations within authorised evidence.	Evidence access and derivative visibility.
Entity	Find identifiers/relationships/related evidence for one entity.	Canonical entity permissions.
Product	Find products by case/type/status/version/recipient metadata.	Product/dissemination permissions.

17.2 Search result architecture

- Result type SHALL be explicit.
- Primary label plus disambiguating metadata SHOULD be present.
- Relevant status/classification SHALL be visible when authorized.
- Search snippets SHALL not expose restricted content.
- Results SHALL resolve to canonical objects, not stale search copies.
- Case-context search MAY show case relevance alongside canonical object identity.
- Every result carries a canonical back-link (API resource and the UI route in a readable case context); an extract result opens its canonical extract resource (/evidence-extracts/{extractId}) within its evidence item. Claims are labelled as claims, never as facts; relevance is labelled as text match, never as risk. [v0.1.3 · CR-I5-01, CR-I5-03, CR-I5-07]
- Derived-text results (object type DERIVED_TEXT, marker DERIVED) are labelled machine-derived ("not the original evidence") with their review state, and link to the derived text and to the original evidence page. Machine translations are shown side by side with their source text. [v0.1.4 · CR-N-11, CR-N-10]

17.3 Facets

Object family	Minimum useful facets
Cases	status, owner, classification, jurisdiction, date
Entities	type, jurisdiction, status, identifier type, linked case
Evidence	source type, file type, date, integrity/verification state, case
Relationships	relationship type, status, valid time, case
Events	event type, date range, entity/asset
Value flows	class, mechanism, currency, date, origin/destination
Products	type, status, reviewer, classification, date

18. Cross-Object Discovery and Backlinks

Every canonical object SHOULD expose meaningful inbound and outbound associations, subject to permissions. Backlinks reduce the need for analysts to remember where an object was used.

Object	High-value related information
Evidence	source, extracts, facts/claims, entities mentioned, relationships supported, assessments/products citing it
Entity	identifiers, relationships, assets, events, value flows, evidence, case contexts
Relationship	two endpoints, evidence, events, case contexts, graph views
ValueFlow	origin/destination, legs, evidence, events, related assets/contracts, hypotheses
Hypothesis	supporting/contradicting evidence, indicators, gaps, assessment
Assessment	hypotheses, evidence chain, gaps, reviews, product versions

19. Deep Linking and Canonical URLs

Information objects that users need to cite, review or share internally SHOULD have stable deep links. URL design is implementation-specific, but semantics SHALL distinguish canonical object identity from case context.

Conceptual pattern:
/objects/entities/{entity_id}

```

/cases/{case_id}/entities/{entity_id} # contextual view or wrapper
/evidence/{evidence_id}#page=12      # precise citation where supported
/products/{product_id}/versions/{version_id}

```

- Deep links SHALL re-check authorization at access time.
- A case-context URL SHALL not create a second canonical entity.
- Precise evidence anchors SHOULD preserve page/region/citation location when technically possible.
- Links in exported intelligence products SHOULD use stable identifiers even when interactive URLs are not externally available.

20. Information Scent and Context Cues

Users should understand what an item is and why it is relevant before opening it. Information scent SHALL rely on concise metadata, not sensational visual emphasis.

Surface	Recommended scent
Entity result	type, canonical name, jurisdiction/identifier, status, relevant case count if permitted
Evidence result	title/file, source, date, case relevance, verification state
Relationship	relationship label, endpoints, time, evidence count/status
Value flow	origin → destination, class, value/date if known, evidence count
Hypothesis	statement, status, support/contradiction counts, open gaps
Product	type, case, version, status, classification

21. Permission-Aware Information Architecture

Permissions shape what information exists from the user's perspective. IA SHALL not rely on disabled links or empty sections to reveal restricted objects.

- Restricted objects MAY be completely absent from navigation/search when policy requires non-disclosure.
- If the user may know an object exists but not see details, the system SHOULD show a policy-safe placeholder rather than misleading "not found".
- Counts, facets and relationship degree SHALL be permission-aware so they do not leak hidden objects.
- Protected-source identity SHALL not appear in general entity indexes, autocomplete, search snippets, graph labels or exported object lists.
- Cross-case backlinks SHALL only list cases the user is allowed to know exist.

22. Derived Views: Graph, Timeline and Dashboards

Derived views are navigation and analytical projections. They SHALL remain anchored to canonical objects and respect the same permissions.

Derived view	IA role	Canonical return path
Investigation Graph	Explore network topology and evidence-backed relationships.	Node → canonical entity/asset; edge → relationship/value-flow record.
Timeline	Explore chronology and temporal clustering.	Item → canonical Event or related object.
Value-flow diagram	Explore movement/transformation of value.	Leg/edge → ValueFlow record and evidence.
Dashboard/ summary	Orient user to work state and notable counts.	Metric/list item → underlying canonical collection.

Derived-view safeguard

No graph node size, dashboard ranking, timeline proximity or search relevance score SHALL become an implicit semantic category such as “most suspicious”. Analytical scoring, if later introduced, requires explicit separate specification and explanation.

23. Content Lifecycle, Versioning and History

Information type	Lifecycle IA requirement
Evidence original	Immutable identity; new file creates new object/version rather than overwrite.
Entity	Canonical identity persists through edits; merge/unmerge/supersession history remains discoverable.
Controlled vocabulary	Stable term IDs; deprecated terms remain interpretable.
Assessment	Immutable published/final versions; later revision is a new version.
Intelligence product	Version chain with current/superseded/retracted states.
Dissemination	Immutable record of exact shared version and recipient context.

24. IA Requirements Catalogue

ID	Normative requirement
IAR-001	Canonical reusable objects SHALL have stable identifiers independent of case context.
IAR-002	Case workspace SHALL link to canonical objects rather than duplicate them.
IAR-003	Global navigation SHALL expose information domains, not backend service/module names.
IAR-004	Graph, timeline and flow diagrams SHALL be derived views with canonical return paths.
IAR-005	Relationship objects SHALL be independently addressable and provenance-bearing.
IAR-006	Search SHALL be permission-aware at result, snippet, count and facet levels.
IAR-007	Protected-source identities SHALL be excluded from general indexes and navigation.
IAR-008	Labels SHALL preserve analytical uncertainty and avoid accusatory structural terminology.
IAR-009	Controlled vocabularies SHALL use stable IDs and retain historical semantics.
IAR-010	Deep links SHALL distinguish canonical identity from case-context views.
IAR-011	Every material assessment SHALL expose navigable links toward supporting evidence chain.
IAR-012	Product/dissemination records SHALL resolve to exact immutable product version.
IAR-013	Unknown values SHALL be represented distinctly from none/zero/false.
IAR-014	Cross-case links SHALL not imply access or reveal restricted case existence.
IAR-015	Entity, relationship and value-flow status/class SHALL be available as structured metadata for search/filter/browse.
IAR-016	Historical/superseded records SHALL remain discoverable to authorized reviewers without being confused with current state.
IAR-017	Case-context relevance SHALL be presented separately from canonical object properties.
IAR-018	IA changes to canonical labels/taxonomies SHALL be version-controlled and reviewed for migration impact.
IAR-019	Browse and search paths SHALL converge on the same canonical object.
IAR-020	No derived ranking or visualization attribute SHALL silently become a semantic guilt/risk classification.

25. IA-to-UX Handoff Contract

Decision area	IA owns	UX owns	Joint acceptance
Navigation	Destinations, hierarchy, route semantics, information grouping.	Interaction of menus/tabs, responsive behaviour, focus/keyboard state.	Users can predict where information lives and move without losing task context.
Labels	Canonical destination/object labels and taxonomy terms.	Microcopy, action wording, validation/help text.	Terminology is consistent and comprehensible.
Object pages	Information hierarchy and related-object sections.	Interaction patterns, editing behaviour, feedback/error states.	Canonical identity and provenance remain clear during tasks.
Search	Scopes, facets, result metadata, permission-aware findability.	Search input interaction, query feedback, empty/loading/error states.	Authorised users can find objects without leakage.
Graph/timeline	What canonical information each projection represents.	Manipulation, selection, zoom/filter interactions and accessibility.	Projection never becomes detached from canonical semantics.

26. IA Validation and Testing

26.1 Validation methods

- Tree testing for primary destinations and case workspace hierarchy.
- Card sorting for labels and controlled-vocabulary groupings where domain language is uncertain.
- Findability tasks across case, entity, evidence and product scenarios.
- Permission leakage tests for hidden objects, counts, facets and backlinks.
- Cross-object traceability tests from product → assessment → hypothesis → evidence → source.
- Deep-link tests for canonical versus case-context identity.
- Terminology consistency audit against Data Model, Typology Catalogue and UX Specification.

26.2 MVP IA acceptance scenarios

1. An investigator can locate a case, find its evidence, open a cited entity, inspect a relationship and return to the case without encountering duplicate object identities.
2. A reviewer can start from an intelligence product and navigate backward to assessment, hypothesis, evidence extract and original evidence/source.
3. A data steward can find an entity from global search, identify conflicting identifiers, inspect merge history and see authorised case contexts.
4. A user without access to a sensitive case cannot infer that case through search counts, backlinks or graph degree.
5. A protected-source identity is absent from global search and ordinary case navigation while source-derived evidence remains usable.

27. IA Definition of Done


- Canonical object and case-context responsibilities are documented.
- Primary navigation and case workspace hierarchy are defined.
- Object taxonomy and controlled-vocabulary boundaries are defined.
- Search scopes, facets and result information scent are specified.
- Cross-object links/backlinks and deep-link semantics are specified.
- Permission-aware findability and protected-source exclusions are verified.

- Terminology matches the canonical data model and UX Specification.
- Representative tree/findability tests pass for MVP workflows (target; no such tests have been run at this baseline). *[v0.1.1 · A01]*
- No IA decision converts analytical status or derived ranking into implicit guilt/risk classification.

28. Traceability to CS-AML Suite

IA area	Primary upstream references
Canonical objects	Data Model Specification; SRS DR requirements.
Case workspace	PRD MVP journey; Investigation Methodology; MVP Engineering E2.
Evidence hierarchy	Data Model; Control Guide EVD/SRC controls; UX Evidence section.
Entity/network hierarchy	Data Model; Control Guide ENT/REL/AST; MVP Engineering E4.
Value-flow structure	Data Model; Methodology; UX Follow-the-Value; E5.
Analytical reasoning	Typology Catalogue; Investigation Methodology; UX analytical sections; E6.
Search/graph	SRS SCH/GRF; Technology Architecture; E7.
Products/dissemination	Methodology; Control Guide QUA/DIS; UX review/dissemination; E8.
Permissions/findability	Technology Architecture; Control Guide SEC; UX permission handling; E1.

Annex A — Reference Navigation Tree



This tree is the MVP IA baseline, not a pixel-level navigation design. UX determines interaction presentation; engineering may implement routes differently as long as semantic hierarchy and canonical identity rules remain conformant.

Annex B — Core Label Dictionary

Canonical label	Meaning / usage
Case	Investigation context with question, scope, roles and analytical work.
Source	Origin of information or evidence.
Evidence	Preserved material used to support or challenge propositions.
Evidence Extract	Precise excerpt/location linked to parent evidence.

Canonical label	Meaning / usage
Entity	Reusable person/organisation/account/address/etc. identity object.
Relationship	Evidence-bearing typed connection between objects.
Asset	Value-bearing object such as property, shares, vehicle, vessel or crypto asset.
Event	Time-bounded occurrence linked to evidence.
Value Flow	Movement/transformation of economic value with explicit evidentiary class.
Indicator	Observed fact/pattern relevant to analysis; not proof.
Typology	Reference pattern/mechanism from the CS-AML catalogue.
Hypothesis	Testable explanation.
Intelligence Gap	Material unknown that could change assessment.
Assessment	Analyst judgement with confidence, basis, alternatives and gaps.
Claim	Attributed assertion from a source; never shown as a fact without a recorded verification decision. [v0.1.1 · A10]
Fact	Proposition, separate from any claim, created with references to supporting evidence (mandatory) and, optionally, supporting claims, together with its recorded creation decision [v0.1.1 · C02]; status Provisional, Established, Disputed or Superseded. [v0.1.1 · A10]
Confidence	High, Moderate, Low, Insufficient basis (wire HIGH, MODERATE, LOW, INSUFFICIENT_BASIS). Insufficient basis is not a level below Low and is never shown as Low, blank or zero. [v0.1.1 · A09]
Classification	Public, Internal, Sensitive, Restricted, Source-protected (wire PUBLIC, INTERNAL, SENSITIVE, RESTRICTED, SOURCE_PROTECTED); access labels are additional restrictions. [v0.1.1 · A08]
Intelligence Product	Versioned analytical output intended for review/use/dissemination.
Dissemination	Controlled record of external/internal sharing of a specific product version.

Annex C — Information Architecture Change Record

Field	Description
Change ID	Stable IA change identifier.
Affected labels/objects	Navigation, taxonomy, labels or hierarchy affected.
Reason	User research, product change, data-model change, legal/security requirement.
Migration impact	Routes, saved links, taxonomy mappings, historical labels, search facets.
UX impact	Interaction or usability implications requiring UX review.
Data impact	Canonical schema or vocabulary impact requiring data review.
Security impact	Permission/findability implications.
Decision / approver	Approved, rejected or deferred with owner/date.
Effective version	IA release in which the change applies.